

RECON 6

This is a Python translation and upgrade of RECON5-5, a Fortran program (©Sukumar & Breneman, 2001) that computes **TAE** (Transferable Atom Equivalent) molecular descriptors — quantum-chemistry-derived property descriptors assembled from a precomputed atom-type database, without running a new quantum calculation for each molecule.

Transferable Atom Equivalent Reconstruction (**TAE-RECON**) is a computational chemistry method that rapidly generates highly accurate quantum mechanical molecular descriptors. Originally developed by the late Prof. Curt Breneman's group at RPI, it drastically accelerates molecular modeling by assembling pre-computed atomic electron-density fragments, skipping time-consuming *ab initio* calculations.

The core of the algorithm is the Atomtyper and TAE descriptor pipeline: reading a molecule, classifying each atom into a 49-character alphanumeric code, matching that code against the TAE database, and combining the matched atomic descriptors into whole-molecule descriptors. This includes several features not present in earlier versions, and has been prepared with the help of Claude (Sonnet 4.6)

How TAE-RECON Works

- **Fragment-Based:** It uses Bader's Quantum Theory of Atoms in Molecules (**AIM**) to create a library of atomic charge density fragments, whose properties are computed at quantum mechanical accuracy.
- **Reconstruction:** When given a large molecule, or a large database of molecules, the **RECON** program matches atom types, retrieves the appropriate precomputed density fragments and stitches them together to reconstruct the molecular electron density and molecular properties.
- **Efficiency:** It calculates physical properties (like electrostatic potentials, local average ionization energies, Fukui functions, electron density Laplacians, etc) for tens of thousands of molecules in seconds.

Key Applications

- **QSAR/QSPR Modeling:** TAE-RECON has been used extensively in drug and materials discovery and toxicology to build highly predictive models for molecular properties, and for protein separations.
- **Database Mining:** Rapidly screens large chemical databases or virtual compounds for promising drug candidates.

Supported input formats

- SDF (MDL molfile, V2000)
- MOL2 (Tripos SYBYL)
- PDB (via CONECT records, or distance-based connectivity using a bond-length table)
- Gaussian `.com`/`.gjf` input, Cartesian coordinates only (no z-matrix/internal-coordinate support). Connectivity is always distance-based, the same as a CONECT-less PDB file, since the format has no bond list at all.

- ORCA input (.inp`/orca`), Cartesian `\* xyz <charge> <multiplicity>` coordinate block only (other ORCA input styles, e.g. internal coordinates, are not supported). Everything before the `\* xyz` marker (route lines, `%block ... end` sections) is skipped automatically. Connectivity is distance-based, same as Gaussian/PDB.
- SMILES (a simplified parser matching the capabilities of the original

`rsmiles.f` — see **Known limitations** below)

Installation

```
```bash
```

```
pip install -e .
```

```
```
```

Requires Python 3.8+. No third-party dependencies — the standard library is sufficient.

Usage

- **Command line** (recommended)

```
```bash (or Windows Terminal)
```

```
python -m recon6 --data-dir /path/to/DATA molecule.sdf -o results.csv
```

```
```
```

```
```bash (or Windows Terminal)
```

```
python -m recon6 --data-dir /path/to/DATA --fmt smiles smiles_list.txt -o results.csv
```

```
```
```

The bond-length table is read automatically from ``<data-dir>/bond`;`

pass ``--bond-file`` only if you need to override that location.

Run ``python -m recon6 --help`` for all options.

- **As a library**

```
```python
```

```
from recon6.recon import ReconConfig, run_recon
```

```
config = ReconConfig(
```

```
 data_dir="/path/to/DATA", # directory of TAE .dat files
```

```
 input_files=["molecule.sdf"],
```

```
 fmt="sdf", # or "mol2", "pdb", "gaussian", "orca", "smiles", "auto"
```

```
 output_csv="results.csv", # optional
```

```
 # bond_file defaults to '<data_dir>/bond'; pass an explicit path here only if your bond-length
 # table lives somewhere else.
```

```
)
```

```

results = run_recon(config)

for r in results:

 print(r["Molecule"], r["Energy"], r["chi"])
 ...

```

Each entry in `results` is a dict of descriptor name -> value (Energy, Population, VOLTAE, SurfArea, SIK, SIG, SIEP, Fuk, Lapl, chi, and the associated min/max/bin-fraction variants). The default CLI output filename is `recon.csv`.

### Distance-based connectivity and bond-order inference

PDB, Gaussian, and ORCA input share a common connectivity module (`recon6/connectivity.py`) for the case where no explicit bond list is available: bonds are inferred from atom-pair distances against a reference bond-length table, then a bond-length-ratio heuristic flags likely double bonds for terminal (single-heavy-neighbor) N/O atoms and sulfonyl/sulfate sulfur centers - the only cases where bond order actually affects standard-valence hydrogen counting (see "Hydrogen addition" below). This was validated against real PDB ligand structures: a measurably shortened bond (e.g.  $\sim 0.85\text{--}0.89\times$  the single-bond reference length for C=O) reliably distinguishes a real carbonyl/imine from an ordinary single bond, while a borderline case (a genuine C-OH at  $\sim 0.98\times$ ) is correctly left alone. This inference does not attempt to detect aromatic-ring unsaturation (uniform bond lengths carry no such signal) - it is deliberately narrow in scope, covering only the well-separated cases that were directly validated.

For Gaussian and ORCA input, this bond-order inference is computed (it's part of the shared connectivity logic) but has no opportunity to matter in practice for hydrogen-counting purposes, since neither reader offers hydrogen addition at all - both formats conventionally specify every atom including hydrogens explicitly as part of a real quantum-chemistry input, and "completing" such a file by guessing hydrogen positions would misrepresent what was actually submitted to those programs. See "Hydrogen addition for H-less input" below.

### Preserving response/property data and row alignment

Many real workflows attach a modeled response (e.g. an assay result) or other metadata to each SDF molecule as an `` data field, and need that data lined up with RECON's computed descriptors afterward for machine learning.

This instead `` fields directly while reading the SDF (`readers/sdf.py`, see the `data\_fields` dict on each parsed molecule) and carries them through alongside each molecule's descriptors by tracking input position internally. If a molecule fails to parse or fails during descriptor computation, it is dropped from `recon.csv` entirely (no blank row, since most ML pipelines treat a half-populated row as noise or a hard error) - its row simply does not appear, and its data fields are dropped with it. Because every surviving row's data fields are read directly from that same molecule's SDF block rather than matched up afterward by name, the response/property columns can never become misaligned with the molecule they belong to, regardless of which molecules earlier in the file were skipped.

This is on by default (`ReconConfig.include\_data\_fields=True`, CLI `--no-data-fields` to disable). The extra columns appear after the descriptor columns, named after their SDF tag (e.g. `LIC50`, `EXTREG`).

## Included data

This package includes the TAE atom-type database (900+ binary `.dat` files) and a bond-length table — these must be supplied via `--data-dir` (or `ReconConfig.data_dir`).

The bond-length table is expected at `<data-dir>/bond` by default; use `--bond-file` (or `ReconConfig.bond_file`) only if yours lives elsewhere.

## Package layout

...

recon6/

- fortranio.py    Reader for gfortran unformatted sequential files
- taedat.py     Binary TAE .dat record parser
- periodic.py    Element symbol / atomic number table
- element.py    Atom-name -> element symbol normalizer (ele.f)
- ringid.py     Ring-membership detection (ringid.f)
- atyper.py     49-character atom-type code generator (atyper)
- gettae.py     TAE database index + atom-type matcher (gettae)
- bonds.py      Bond-length table loader
- descriptors.py   Per-molecule descriptor accumulation (qmf)
- sparse.py     Sparse bond-order matrix (memory optimization)
- connectivity.py   Shared distance-based connectivity + bond-order inference (PDB, Gaussian, ORCA)
- hydrogenate.py   Standard-valence H-addition for H-less SDF/MOL2/PDB input
- recon.py      Orchestrator: ReconConfig, run\_recon
- \_\_main\_\_.py   CLI entry point
- readers/
  - sdf.py       SDF (MDL V2000) reader
  - mol2.py      MOL2 (Tripos) reader
  - pdb.py       PDB reader (CONNECT or distance-based)
  - gaussian.py   Gaussian .com/.gjf reader (Cartesian coordinates only)
  - orca.py      ORCA input reader (Cartesian '\* xyz' block only)
  - smiles.py    Simplified SMILES parser

tests/

- test\_readers.py      Unit tests for readers, ele, periodic

```

test_rigid.py Ring detection unit tests
test_hydrogenate.py H-addition geometry and formula-consistency tests
test_pdb_reader.py PDB element-column, deuterium, bond-order inference tests
test_gaussian_reader.py Gaussian .com/.gjf parsing tests
test_orca_reader.py ORCA '* xyz' block parsing tests
test_data_fields.py SDF data-field extraction, CSV quoting, alignment tests
test_integration.py 974-molecule SDF batch vs. Fortran reference
test_integration_toxx.py 278-molecule SDF batch vs. Fortran reference
test_integration_smiles.py 106-line SMILES batch vs. Fortran reference

```

```

` ` `

```

## Running the tests

```

` ` ` bash

```

```

python -m unittest discover -s tests

```

```

` ` `

```

Most of the integration tests depend on sample data (the TAE DATA directory and several large molecule files used during development). Tests whose data isn't present are skipped rather than failed. To point the suite at your own copies of this data, set any of the following environment variables before running (see `tests/test_config.py` for the full list and defaults):

```

` RECON6_DATA_DIR`, ` RECON6_GAC_SDF`,
` RECON6_GAC_FF`, ` RECON6_TOXX_SDF`, ` RECON6_TOXX_FF`,
` RECON6_TOXX_MOL2`, ` RECON6_BENZOX_TXT`, ` RECON6_BENZOX_FF`,
` RECON6_PDB_NNA`, ` RECON6_PDB_SV6`.

```

## Validation

This translation was validated against real Fortran RECON5 output on three independent test sets:

Test set	Molecules	Format	Result
GAC_withH.sdf	974	SDF	Despite the filename, 971/974 of these molecules have

**\*\*zero\*\*** explicit hydrogens in the source file, so the Fortran reference for those rows reflects its H-less, unsaturated behavior - not a meaningful comparison now that this release adds missing H's. The 3 molecules that genuinely had explicit H's match the Fortran reference to <0.05%; H-addition for the other 971 is validated by formula consistency (added H count matches each molecule's formula) instead. 3 further molecules are excluded for an unrelated, pre-existing input defect (disconnected halogen atoms) - see Known limitations. |

| tox2.sdf | 278 | SDF | 278/278 match to <0.05% (this file is properly H-saturated already) |

| Benzoxazines.txt | 106 lines | SMILES | 94 lines use syntax supported by the original `rsmiles.f`; of those, 90+ match Fortran energies exactly |

Tolerances of <0.05% reflect ordinary floating-point/formatting differences (the Fortran reference output is truncated to ~5-6 significant figures in its text format).

Additional validation beyond the table above: the MOL2 reader was checked against a 278-molecule MOL2 conversion of the same tox2 compound set (reordered relative to the SDF version, so validated by comparing the multiset of computed Energy values rather than row order - 278/278 matched). The PDB reader, hydrogen-addition, and bond-order inference were validated against two real PDB ligand extracts (49 and 103 heavy atoms, no CONECT records, no hydrogens, one with partial deuterium labeling) - added hydrogen counts matched each ligand's molecular formula once carbonyl/sulfonyl bond-order inference was in place. The Gaussian and ORCA readers were cross-checked against each other (identical methane input via both formats produces identical descriptors) and against the same bond-order-inference logic validated for PDB.

The package has also been run successfully (no crashes) on much larger external datasets, including a 51,449-molecule SDF batch and several other independent SDF datasets supplied during development.

### Known limitations

**SMILES parser:** The original `rsmiles.f` is a minimal SMILES reader supporting only plain atoms, single/double/triple bonds, branches, and single-digit ring closures — it has no concept of stereochemistry, formal charges, or bond-direction notation. This release matches that scope and explicitly **rejects** (rather than silently mis-parsing) SMILES containing:

- Bracket atoms, e.g. `[C@H]`, `[N+]`, `[O-]`
- Bond-direction markers `/` and `\`
- A leading `(` before the first atom

A small number of SMILES with reused ring-closure digits in deeply branched systems may also diverge from the Fortran reference; this affects roughly 4 of 94 parseable molecules in the validation set and is a known, low-priority edge case.

**Disconnected atoms:** If an input structure contains an atom with zero bonds (a data-quality defect — e.g. 3 molecules in the GAC\_withH.sdf test set have chlorine atoms with no bond records at all, despite a nonzero molecular formula), the atom-type matching falls back to an essentially arbitrary TAE database entry. This mirrors the original Fortran's behavior for the same degenerate case (verified by tracing `atyper`'s logic) rather than being a translation bug, but results for such atoms — and the molecule's aggregate descriptors — should not be trusted. A warning is printed to stderr when this is detected.

**Ring sizes:** `ringid` correctly detects rings of size 3 through 6 (plus several fused-ring codes), but the TAE database (`TAE.LIST`) only contains atom-type entries for 5- and 6-membered rings, so 3- and 4-membered rings have no matching descriptor data in practice.

**Bond-order inference does not cover aromatic rings:** The bond-length-ratio inference described above (for PDB/Gaussian/ORCA input) only targets terminal N/O double bonds and

sulfonyl/sulfate sulfur, which were directly validated. Aromatic ring bonds have essentially uniform length regardless of formal bond order, so there is no geometric signal to detect their unsaturation this way. In practice this only matters for PDB input, since Gaussian and ORCA never get automatic hydrogen addition (see above) - a heavy-atom-only PDB structure with an aromatic ring, run through the H-adder, could miscount hydrogens on that ring for the same reason. If this matters for your workflow, treat it as a follow-up rather than an assumption this release resolves silently.

### Hydrogen addition for H-less input

Many real-world SDF/MOL2 files only specify heavy atoms, and PDB ligand extracts very commonly lack hydrogens entirely. The original Fortran simply logged ``WARNING: Molecule contains no Hydrogens`` for any molecule with more than 10 atoms and zero explicit hydrogens, then proceeded with that incomplete connectivity. This release instead detects that condition (``recon6.hydrogenate.needs_hydrogens``, matching the same ">10 atoms, 0 H" trigger) and saturates missing valences using standard valence rules (the same default valences RDKit/OpenBabel apply), with new hydrogen positions placed via local-geometry heuristics (linear / trigonal planar / tetrahedral, depending on each atom's existing neighbor count) rather than the Fortran's original no-op.

This is on by default for SDF, MOL2, and PDB input; disable it with ``ReconConfig(auto_add_h=False)`` or the CLI's ``--no-add-h`` flag. SMILES input already gets H-saturated by the SMILES parser itself (matching the original ``rsmiles.f``), so the flag has no effect there. For PDB specifically, the same standard-valence formula needs to know real bond order to avoid overcounting H on double-bonded atoms (e.g. a carbonyl oxygen) - see "Distance-based connectivity and bond-order inference" above for how that's handled when the input format itself carries no bond-order information.

Gaussian and ORCA input never get automatic hydrogen addition: Both formats conventionally specify every atom including hydrogens explicitly as part of a real quantum-chemistry input – silently "completing" a heavy-atom-only file by guessing hydrogen positions would misrepresent what was actually submitted to those programs, so this is not offered as an option at all for ``read_gaussian_com`` / ``read_orca_xyz`` (there is no ``auto_add_h`` parameter on either). A heavy-atom-only Gaussian or ORCA file is read and processed exactly as given, with no hydrogens added.

Important caveat - partially-specified hydrogens: The trigger only fires when a molecule has **zero** explicit hydrogens. If a structure has *some* but fewer than its formula implies (for example, a deuterium-labeled compound where only the non-deuterated hydrogens were given explicit positions), no automatic addition happens, and the molecule is processed as-is. This is intentional: a low-but-nonzero H count could reflect deliberate isotope labeling or another legitimate partial specification, and guessing which atoms are "missing" their hydrogens in that case is a different, harder problem than full saturation from zero. If you need this case handled, treat it as a follow-up request rather than an assumption this release makes silently.

When hydrogens are added, a one-line note is printed to stderr per molecule (``Note [name]: added N hydrogen(s) ...``) so it's visible in batch logs.

## Authors

Several authors have put in a LOT of work developing the original package over the years (in alphabetical order): Curt Breneman, William P. Katt, Martin Martinov, Marlon Rhem, Dominic Ryan, N. Sukumar, Tracy R. Thompson, Christopher Whitehead, Dechuan Zhuang.

## References

1. N. Sukumar and Curt M. Breneman, "*QTAIM in Drug Discovery and Protein Modeling*" in "The Quantum Theory of Atoms in Molecules: From Solid State to DNA and Drug Design" C.F. Matta & R.J. Boyd, Eds. (Wiley-VCH, 2007) ISBN: 9783527307487
2. Christopher E. Whitehead, Curt M. Breneman, N. Sukumar and M. D. Ryan, "Transferable Atom Equivalent Multi-Centered Multipole Expansion Method" *J. Comp. Chem.* **24**, 512-529 (2003) DOI: 10.1002/jcc.10240
3. Breneman, C.M., et al., *Electron Density Modeling of Large Systems Using the Transferable Atom Equivalent Method*. Computers & Chemistry, 1995. **19**(3): p.161.
4. Thompson, T.R., *Construction of a Library of Transferable Atom Equivalents*, in *Chemistry*. 1994, Rensselaer Polytechnic Institute: Troy.
5. Bader, R.F.W., *Atoms in Molecules: A Quantum Theory*. 1990, Oxford: Oxford Univ. Press.